AgentVQA

Human evaluation · Participant instructions

Protocol: agentvqa-human-calibrated-v1
Study address: https://advaitgupta.github.io/human_eval1/
Contact: advaitgupta@ymail.com

1. What you will do

Answer multiple-choice questions about websites and apps, games, egocentric videos, robotics, and spatial planning. Choose the single best option using only the supplied visual material, question, episode goal, and action history. Some questions have two choices; others have more than four. Read every available option.

You may choose 100 or 200 questions. Both sets are fixed across visits and participants. The 100-question set is a subset of the 200-question set. Each domain contributes 20 or 40 questions, respectively, distributed equally among its source datasets. You may take breaks, skip questions, revisit answers, or end at any time. This is the fixed question set used for the calibrated original-option comparison; it is separate from the earlier human-evaluation study.

2. Work independently

Please do not use AI assistants, search engines, external documentation, the released dataset, answer keys, or another person's answers. Do not look up the study examples. General knowledge is fine. If you previously saw or helped curate these questions, record that on the start page.

No correctness feedback or model predictions are shown. There is no time limit. Select your best answer even if uncertain; if you cannot judge an item, leave it unanswered and explain why in the optional comment. Use the flag checkbox for ambiguity, missing information, or media problems. A flag is not an answer.

3. Read all supplied context

Episode Goal states the task the agent is trying to complete. Action History lists actions before the current observation. These fields appear only when supplied by the dataset. Game descriptions, reward rules, spatial instructions, and trajectory notes inside a question are part of the task and should be read.

Scroll through the page to check all supplied images, the full question, episode goal, action history, and every option. Some content may be below the visible part of the page. Confirm that you are inspecting the correct picture and its printed labels before choosing an answer; do not assume the first visible picture is the only one.

The selectable Options list is authoritative if an older choice list appears inside the question. Choose its letter, not an embedded state number or another labeling scheme. For example, option A may refer to Image 0. Do not assume an unshown action or invent information.

4. Images, coordinates, and precise markers

Each image starts fitted to its viewing box so the complete image is visible without scrolling inside the image. The original pixel dimensions are retained. Use +, −, Fit, or 100% to change the zoom. After zooming in, you may need to drag or scroll horizontally and vertically to inspect the rest of the image. Fit restores the complete-image view; Expand opens a larger viewer. For image-point grounding options, values are [x, y] in source-image pixels with the origin at the top-left; x increases to the right and y downward.

Every image-point grounding option has a neutral letter bubble and a small black dot at the exact coordinate. The center of the black dot—not the bubble, its edge, or a connecting line—is the candidate point. Letters may be offset to avoid collisions, but the dots are never moved. Zoom in to inspect a boundary. Hide markers to inspect the original image, or display one option's markers at a time. Marker colors never indicate the correct answer.

Warning: some choices are hard negatives. A point can be extremely close to the correct object or button yet still fall just outside the required target region. Being nearby is not sufficient. Zoom in and check the exact black-dot center against the object's boundary; compare all candidate points carefully before answering.

An option without an image-point coordinate, such as a button press or None of the above, has no added point. An out-of-image coordinate is not clamped to the image edge. Existing arrows, numbered points, grids, and trajectories in an image remain unchanged. Symbolic points such as p_7, grid cells such as c_2, and game-board coordinates are not source-image pixels; use the labels already shown in the image.

5. Multiple images and videos

Multiple images appear in their original dataset order and are labeled Image 1, Image 2, and so on. Read any original labels printed inside an image: they may use different numbering. No panels are removed. Video questions provide the complete clip. You may play, pause, replay, seek, change playback speed, or use the small-step controls to inspect events.

Zoom and full-video access are allowed in this human study. They are viewing aids, not a claim that every evaluated model used identical visual preprocessing. If media fails to load, try reloading it or opening the media link. Flag persistent problems and leave a comment rather than guessing from a missing image or clip.

6. Navigation, comments, and saving

Use domain tabs, the dataset selector, question-number buttons, or Previous and Next to navigate. Shaded number buttons have an answer; a small amber dot means you flagged the item. Selecting an option does not automatically advance. You can clear or change an answer at any time before your final export. Comments or reasoning are optional and saved with that question.

Your progress is saved in this browser on this device. Closing the page and returning to the same address lets you resume, unless browser storage was cleared or unavailable. Switching to 200 questions retains answers from the 100-question set. Download a JSON backup before changing devices or clearing storage; use Import JSON to restore it. Avoid editing the same session in two tabs simultaneously.

7. End, download, and send your responses

Choose End & export, then Generate PDF + JSON. The PDF contains the selected questions, all available choices, source IDs, context, media references and image previews, your chosen answers, and comments. The JSON additionally preserves exact source fields, media identities, study version, flags, and viewing-interaction records. Unanswered items remain explicitly unanswered. No correct answers are included.

Your browser may require permission for multiple downloads. Separate download buttons remain available if either file is blocked. The study is not submitted by downloading. Use Open email, attach the downloaded JSON file manually, and send it to advaitgupta@ymail.com. You may also attach the PDF. Where supported, Share JSON opens your device's file-sharing menu; choose email and verify the recipient. The page cannot automatically attach a file to a mailto message or confirm that your email was sent.

8. Privacy and participation

Use a participant code or alias; do not enter passwords, tokens, or sensitive personal information. Participation is voluntary. Responses, comments, approximate visible time per question, and counts of viewing actions are stored locally and included in your exports. They are not automatically uploaded or emailed by this site.

The page, images, and lossless browser-compatible MarioQA videos are served by GitHub Pages; other videos are loaded from Hugging Face. Those services receive ordinary web requests when you access their content. This interface adds no analytics or tracking service. Your downloaded response files are shared with the researcher only when you choose to send them. You may stop without sending anything. Contact advaitgupta@ymail.com with study questions.

Fixed allocation

Web Agents: 100-question set = 20 questions (5 per dataset); 200-question set = 40 questions (10 per dataset).
Datasets: AitW, MONDAY, Screenspot, Screenspot-Pro.

Games: 100-question set = 20 questions (5 per dataset); 200-question set = 40 questions (10 per dataset).
Datasets: Atari, GameQA, MarioQA, VideoGameBunny.

Egocentric Videos: 100-question set = 20 questions (4 per dataset); 200-question set = 40 questions (8 per dataset).
Datasets: EgoPlan-Bench, OpenEQA, Perception-Test, VidEgoThink, VSI-Bench.

Robotics: 100-question set = 20 questions (4 per dataset); 200-question set = 40 questions (8 per dataset).
Datasets: ERQA, ManipBench, Robo2VLM, Roborefit, Share-Robot.

Spatial Planning: 100-question set = 20 questions (4 per dataset); 200-question set = 40 questions (8 per dataset).
Datasets: 3DSRBench, EmbSpatial-Bench, OmniSpatial, SpaCE-10, Spatial-MM.

Dataset revision: 29224a4605b0020fb5fca35a8ec93b6300699c85
Evaluator revision: calibrated-original-option-cohort
The complete dataset release, source fields, and media identities are recorded in response JSON files.