AgentVQA · Open-ended review

Review instructions · https://advaitgupta.github.io/open-ended/

What you are reviewing

Judge saved open-ended answers for 200 questions across five domains and 23 datasets. Each question has outputs from seven models. No new model calls are made by this page. Choose all models or one model to review; the question set stays fixed.

Read the task first

Read the question, Episode Goal and Action History when provided. Inspect every image or the video. The exact prompt and output-format instruction given to the selected model are available below its answer. The original multiple-choice record is optional reference only: it may have different wording, and “None of the above” is not an open-ended target.

How to judge

Correct: the answer satisfies the task. Incorrect: the answer is substantively wrong. Unsure: you cannot confidently decide. Not gradable: the question, media or response prevents a defensible judgment. Leave unreviewed answers blank. A formatting warning is not an automatic incorrect grade; record content correctness and explain format or ambiguity issues in your comment.

Points, cells and labels

Numeric point markers use the coordinate convention saved with that model’s prompt, not a guessed convention. The tiny black-dot center is the exact predicted location; the blue bubble may be offset. Out-of-range coordinates are reported, not clamped or silently changed. A board or trajectory cell is outlined as a region. A PICK label highlight locates an existing printed label; use the original image dot as authoritative.

Inspect carefully

Images start fitted in full. Scroll the page to inspect all images and context; after zooming in, pan or scroll within an image. Some grounding targets have close boundaries, so hide the marker and inspect the actual object when needed. Do not accept a point merely because it is close to the target. Video reviewers can pause, seek and replay; the model may have received only the frames listed in its saved protocol.

Independent judgments

Judge the answer against the visual evidence, not against another model’s answer or reputation. Optional anonymous model labels are available. Add a comment for borderline cases, missing information, questionable references or formatting problems. The interface does not automatically score answers.

Saving and exports

Progress autosaves only in this browser. Use an alias rather than personal information. Download JSON backups to move devices; import restores the same catalog’s judgments. End & export offers PDF, JSON and CSV. JSON includes questions, source IDs, model answers, prompts, coordinate overlays, your grades and comments. Review results are not uploaded automatically.

Sending results

Download the JSON and attach it manually to an email to advaitgupta@ymail.com. The email button opens a draft but cannot attach a file or send it automatically. This is a public GitHub Pages interface; only the saved benchmark/model data is hosted, not your private browser annotations.